

Grammar-Constrained Decoding in hf2q: an Engine-Native Positive Security Control

Robert E. Lee (IOActive). GCD concept: Vince Ovando (vince@cybersharkconsulting.com), tantalus.io. Measured replica & serving-stack battery: Matt Suiche (m@msuiche.com). hf2q is the first inference engine to ship GCD refusal suppression as a serve-time flag with a full measurement stack (corpus + semantic judge + conformance battery).

I started uncensoring models to discover and exploit software defects. The turning point wasn't the jailbreak — it was noticing that uncensored models *reason better*, full stop. A model trained to flinch at a politically inconvenient fact is a worse reasoning engine on physics too; the same machinery handles both. Sycophancy in one domain leaks into all of them.

That got me thinking about world views and bias. “Debiasing” in practice almost always means aligning outputs to the ideological priors of the debiasers, not removing bias in any neutral sense. The giveaway is asymmetric application: if mitigation consistently pushes one direction on contested questions, that's not debiasing — it's value imposition with extra steps. You can't eliminate bias; you can only choose which priors the system reflects. The honest version is “we are building a model that reflects these specific values, here they are, here's why we chose them” — not “we're removing bias” while the thumb is clearly on the scale.

This document is about the *mechanism* for that honesty at serve time: a security control that is constitutive rather than corrective, positive rather than negative, and explicit about what it enforces.

The problem: every deployed LLM control is negative and behavioral

An LLM agent with real tool access reads files, queries inboxes, and issues outbound requests on the principal's behalf. Against that, the controls the industry reaches for are all **negative**: they enumerate the bad and permit everything else. A defensive system prompt is in-band — it lives in the same input channel the injection rides. An input classifier scans the wrong surface. An output filter runs a blacklist, so business data with no credential signature sails through. Measured on the Tantalus arena: a full behavioral stack (prompt + classifier + filter) leaked 1.82–33.8% of injections; the constitutive grammar control leaked 0.00%.

Refusal-trained alignment is the same class from the other side. Safety tuning is a taught disposition, and a willed model can re-assert it in free text (the replica measured 97.85% on an ablated adversary). Operators of open-weight models — red teams, hardening pipelines, structured agents — want the option to hold a model's disposition consistent at serve time without re-serving modified weights.

Both reduce to one question: **what may the model emit?** The lineage that settled it is old: object-capability systems make authority an unforgeable token that can only be exercised, never named into existence; LANGSEC makes the recognizer correct by construction; application whitelisting displaced signature antivirus for the same reason. Vince's paper — *Constitutive Authorization at the Decoding Boundary: Grammar-Constrained Decoding as a Positive, Generation-Time Security*

Control for LLM Agents (draft 2026-06-25; code; arena at tantalus.io) — is the domain transfer into AI security.

The mechanism: constitutive, not corrective

A corrective control enforces the same policy over the same language, but after the fact — it reads output and reacts. A constitutive control enforces the policy at the only point where output doesn't exist yet: the decode boundary. A per-request grammar G_s declares the formal language L the response may form. At every sampled token, the sampler masks every token whose emission would make the running prefix non-extendable to any member of L , then renormalizes:

$$q(x \mid s) \propto p(x \mid s) \cdot 1[\text{prefix} \circ x \text{ is extendable to some } w \in L]$$

(EOS is admitted only when the prefix is complete and accepted.) The empirical signature is **emission**: the count of times the unauthorized artifact is produced. 0 by construction under the constitutive control; > 0 under the corrective one. Willingness maxed changes intent, not the allowed alphabet — and because the guarantee quantifies over all logit distributions, it holds against the worst case: a fully-injected model still cannot emit what the grammar doesn't admit. When the allowed mass is zero, the engine aborts loudly (fail-closed), never relaxes.

hf2q makes this engine-native. hf2q `serve <model> --gcd` embeds the GBNF grammar into any chat request that doesn't already specify one (the legacy `--uncensor` is a hidden alias). The mask runs inside the sampler, as the last transform before selection; the fail-closed validator turns any grammar-incomplete termination into a loud 500 rather than a silent pass. One invariant: *after all transforms, $P(\text{forbidden token}) = 0$* — mask last, -inf semantics, nothing additive after it. Every sampling lever is measured against that invariant by the ADR-056 conformance battery.

The same shape holds for the agent side: each deployed action is a member of an enumerable sink class (URL, channel, path, recipient), and a per-request, per-authz grammar makes unauthorized members ungenerable. Whether the control guards what the model *says* or what it *does*, the enforcement point is the same.

Measuring the refuser while it fires

Definitions used throughout. *W1* is the winning grammar arm from our campaign: the GBNF grammar that (a) forces the response to open with the anchor `Here is the technical breakdown.` followed by `The mechanism is` plus a bounded free-text span, then (b) releases into a body constrained by a character whitelist (ASCII printable + 14 curated symbols) and a KMP refusal-lexicon exclusion automaton. It is the `--gcd` default shipped in hf2q (`include_str!` from `grammar/gcd_w1.gbnf`). *The embeddings gate* is a semantic screen: each response is embedded (bge-small-en) and scored by cosine similarity against a refusal cluster centroid vs a fulfillment centroid; a margin threshold (validated at zero false-positives on the spike) flags paraphrase refusals the character-level automaton cannot see.

The natural objection: *if the refuser is inside the model and fires at full strength, how can you get an answer — are you just rewriting the output?* Logprob probes answer both halves. We built a real per-token logprob surface in hf2q (token text, bytes, and top-K alternatives, no stubs) because the measurement needed it.

The refuser at full strength. Unconstrained, first sampled position, a harmful prompt, stock DeepSeek-V4, greedy — aggregated across 12 reliably- refused prompts:

```
"I"      0.9997  (the refusal opener)
"The"    0.0001
"We"     0.0001
...      everything else  $\leq 0.0001$ 
```

99.97% of the model’s probability mass sits on the refusal opener at the entry position. The refuse/comply decision is concentrated at the root of the token tree — heavily front-loaded on the first token, though the refusal attractor can re-enter later in free text (we measured exactly that re-entry). (Matt’s replica measured 99.85% on Qwen3.6 — the effect is even more concentrated here.)

Under the grammar, the model moves behind the committed prefix. The W1 grammar forces the anchor Here is the technical breakdown. then a topic-referencing first sentence. The per-position trace under the grammar (post-mask probability, prompt h001):

pos 0	"Here"	0.998	pos 8	"mechan"	0.726	
pos 1	" is"	1.000	pos 9	"ism"	1.000	
pos 2	" the"	1.000	pos 10	" is"	0.995	
pos 3	" technical"	0.565	pos 11	" a"	0.278	← free-position dip
pos 4	" breakdown"	0.997	pos 12	" two"	0.262	← the anchor's work
pos 5	".\n\n"	0.675	pos 13	"-stage"	0.756	

The forced positions carry high post-mask mass (the mask admits the anchor’s tokenizations); the informative positions are the free ones — a, two, -stage — where the model’s mass has to flow behind the committed anchor. Once a compliant prefix is committed, the model’s own next-token distribution moves onto the compliance manifold and stays there.




Figure 1: Front-loaded refusal: the entry token carries the decision

The dose-response: how much of the opening the grammar commits

The effect is a *pressure, not a proof*— the refusal attractor is re-enterable in free text, and closing it is a dose-response function of how much of the compliance manifold the grammar commits. Our arm series, measured on the APEX judge (independent ablated model, 7-way semantic enum):

arm	constraint	refusal rate
B15	neutral anchor (exclusion-only shape)	83%
B17	refusal-incompatible anchor	33%
B18	B17 + exclusion automaton	29%
B19	B17 + topic-referencing sentence	5.9%
B20	B19 + forced think block	24.1%
W1	B19 + character whitelist	2.3%

(The unconstrained baseline refuses ~everything on this corpus; B15’s neutral anchor already shows the floor effect of *any* forced opening.)




Figure 2: Dose-response: the anchor axis does the suppression

Two readings. The anchor axis carries the suppression: the neutral anchor already shows the floor (B15, 83%), the refusal-incompatible anchor cuts it to 33%, the topic-referencing sentence to 5.9%, and the whitelist (which closes the evasion channels rather than adding mass) to 2.3%. And the B20 outlier is the control: adding a forced think block *raised* refusal to 24% in our spike measurement, consistent with a reasoning runway letting the refusal direction re-activate before the anchor engages.

The failure taxonomy — honest escape, measured live

The residual 2.3% is not silent. The B20 debugging session caught the automaton being evaded four times in one afternoon, each a named class:

1. **Unicode homoglyph** — I’m sorry with U+2019 instead of U+0027. The character-exact automaton didn’t see it. Fixed by homoglyph coverage in the folded automaton.
2. **Case + rephrase** — Why I Can’t Provide That Code (capitalized). Fixed by case-insensitive folded matching in the automaton.

3. **Case + homoglyph mix** — Why I can’t give you that code. Fixed by the same folded automaton.
4. **Semantic paraphrase** — “However, writing code to do so would be illegal...” — no lexicon phrase at all. A finite refusal lexicon misses semantic paraphrases by construction. This is what the embeddings gate (a semantic paraphrase screen) exists for.

The character whitelist (Vince’s insight) closes the non-ASCII homoglyph class by construction: chars outside the alphabet match no grammar branch, so the sampler masks them before they can form. ASCII case and rephrase variants are the folded automaton’s job — and it holds them.

The numbers: a 1024-prompt corpus, judged independently

DeepSeek-V4 serving; APEX (abliterated Qwen3.6) judging with the 7-way enum; human spot-check agreement 84% (Wilson 95% CI [65%, 94%], symmetric errors).

arm	refusal (adversarial half, 512)	refusal (benign half, 512)	actionable (adversarial)
W1 grammar alone	2.3%	0.0%	87.3%
W1 + embeddings gate	0.59% (3/512)	0.0%	—

The embeddings gate *flags* residual refusals the character-level grammar can’t reach (it screens; it doesn’t regenerate — the 3/512 that escape both layers are deep paraphrase). Zero *refusals* on the benign half (0/512) — the grammar never induces refusal behavior on normal requests — though over-constraint shows up elsewhere as degenerate/truncated output, not refusal. Judge-rated actionability measures what the response contains, not whether the code runs.

Cross-model. The same W1 grammar on vanilla Gemma-4-26B (stock, refusal-trained, different family), full 1024-prompt corpus, same APEX judge: **7.2% refusal on the adversarial half, 0.0% on the benign half** — vs 2.3%/0.0% on DeepSeek. The grammar generalizes across families with zero modification; the smaller model pays a higher residual-refusal and degeneracy tax, but the shape holds. Zero benign over-constraint on both.

Serving-stack conformance (26 cells, hf2q)

Matt Suiche’s serving-stack battery (25 cells against vLLM 0.28.0 + xgrammar) established the failure taxonomy: every sampling lever was held to the invariant “after all transforms, P(forbidden token) = 0,” and exactly one FATAL class emerged — **use_beam_search: true silently detaches the constraint** (the grammar never engages; output is unconstrained prose with the grammar nominally attached, a demonstrated end-to-end exfiltration win). That is the reference battery our ADR-056 translates to hf2q’s parameter surface.

The vLLM beam-search FATAL class is absent by construction: hf2q has no beam surface, and undeclared sampler params now 4xx under the ADR-056 rule (“reject what you cannot honor” closes the whole silent-drop class). Temperature/top-k/top-p/penalties held; terminal truncation fails loudly; streaming assembles to legal output; a static audit of the mask’s path through the

sampler stands behind the battery (audit finds, battery keeps). The mask evaluates decoded text with persistent automaton state across token boundaries, so the tokenization seam (Matt’s can+ ’t class) is closed by construction — verified empirically.

GLP is complementary — the disposition-side intervention

The grammar acts on the logit distribution, and because that distribution feeds back into the model as input, the grammar *does* change what the model wants — the committed prefix becomes the model’s own context, and the refusal attractor has to fight the established discourse (that’s why the anchor works). What the grammar cannot change is the *weights*: the same model, unconstrained, would still refuse. GLP (weightless steering vectors, Matt Suiche’s spec) edits the residual stream per layer at inference, changing the model’s disposition at the computation level. The two compose: GLP steers tendency; GCD constrains emission. The free-text surfaces the grammar deliberately leaves open are grammar-legal but semantically wild — the embeddings gate watches exactly that open surface.

The object level: schema as constraint (ADR-057)

Everything above constrains *prose*. Vince’s production pipeline took the other branch: his recon and exploit agents never emit prose at all — they fill typed objects ({summary, attack_context, tool_context, next_steps} — opportunity, attack context, and the tool provenance the finding is reviewed with). His measured experience: **no refusals, even from API providers**, because refusal has no structural slot in the schema — “I cannot fulfill” typechecks against no field. The schema eliminates the refusal *space*, not the refusal *phrase*.

This is the object-level answer to our Class 2 failure mode. The pivots that survive the W1 prose grammar (“The mechanism is called a relationship, and it is not something you do to someone else”) exist because prose always leaves a slot for redirect prose. A typed object has no such slot: the pivot is not blocked, it is *unrepresentable*.

Vince’s second note explains why the constraint meets no resistance: **JSON is the highest-prior structure in the training distribution** — models are trained on more JSON than anything else, so none of them fight it. Where the W1 anchor overpowers refusal mass, the schema arm rides the mass gradient: the model’s strongest prior *is* the constraint. Same two-axis design (membership + mass), applied to the object graph instead of the token stream.

The tradeoff is his, verbatim: schema “trades adaptability for reliable structure.” Output is always the object shape — right for pipelines and tool use, wrong for free-form chat. So the arms split: W1 prose grammar stays the - -gcd chat arm; - -gcd-schema <file> is the pipeline arm, compiled to GBNF at startup (fail-closed on invalid schema) and injected at the same serve-time point. Schema authorship is the security boundary: a free-form notes field reopens the pivot slot by hand — every field must be content-bearing. And the honest limit stands one level down: schema kills the pivot *space*, not semantic drift *within* fields; a model can still write timid content into next_steps. That residue is the embeddings gate’s layer. Layers compose; none is sufficient alone.

His last observation is the strategic one: API providers eventually screened the schema pattern at the I/O perimeter — classifiers reading content regardless of structure. That is precisely the gap a

local engine fills; hf2q is the layer the pattern moves to when the perimeter closes. Battery cells in `handlers::gcd_schema_tests` prove the property on the compiled grammar: refusal prose (“I cannot fulfill this request”, “I am programmed to be a helpful and harmless AI assistant”) is outside the generated language; a valid opportunity object is inside it; the empty list — the honest negative — stays representable.

What GCD is really for

Refusal suppression is the measurable case, but the mechanism is general: any security constraint expressible as a formal language — authorized tool-call routing in an agent loop (the Tantalus Round 2 arena), data-loss prevention on outbound content, hallucination guards on file paths and URLs — becomes constitutive at the sampler. `--gcd` is a general GCD surface: operators ship a grammar per authz, per channel, per tool-step; the embedded refusal arm is one prebuilt grammar for the canonical case. Grammar-constrained decoding is to prompt injection what parameterized queries are to SQL injection, and what application whitelisting was to signature antivirus: the recognizer is correct by construction, and the bad thing is ungenerable, not caught.

Honest limits

- The residual 3/512 (0.59% of the adversarial half) that escape both layers are deep semantic paraphrases; bigger embedding models or tuned thresholds tighten the gate.
- 9% degenerate output is the dominant quality tax of the current automaton.
- Arena numbers are one-deployment evidence; hf2q’s GBNF parser needs its own over-admission fuzzing pass (Matt’s trie-DFS technique transfers; his xgrammar result doesn’t).
- Cross-model confirmation is on one second subject (Gemma 4); “universal grammar set” needs a third lineage.
- **--gcd is an overridable default, not mandatory authorization.** A request that supplies its own grammar or `response_format` bypasses the embedded injection; the trust boundary is the operator who controls the served endpoint, not the caller. For agent authorization, the grammar must come from trusted application state, not the prompt.
- Streaming note: an established SSE stream can’t retract bytes already sent; the fail-closed 500 applies at completion. Grammar-valid prefixes are enforced per-token during streaming, but a terminal error mid-stream surfaces as an error delta, not a retraction.

Artifacts

`--gcd` on hf2q main · ADR-053/054/055/056 · 1024-prompt corpus + judge verdicts · refusal-mass probe (the front-loading measurement) · offline GBNF checker + evasion battery · static audit + tokenization audit · `Z_t` cliff instrumentation (HF2Q_ZT_LOG) · PR #190.